

**One real cohort: honest evaluation is at chance;
split-shopping and test-set tuning climb it to a publishable-looking 0.68**

